

Customer Project Steps

Churn Prediction – Data Science

1. Understand the Business Problem

Customer churn is a critical business metric in the telecom industry. Retaining customers is often more cost-effective than acquiring new ones. This project aims to predict whether a customer is likely to leave the service (churn) based on their usage patterns, complaints, support history, and service preferences.

2. Data Collection and Preparation

We use a synthetic dataset containing 10,000 customer records. Features include:

- Customer Age
- Monthly Charges
- Total Charges
- Contract Type
- Internet Service Type
- Technical Support Availability
- Tenure in Months
- Payment Method
- Streaming Services
- Customer Support Calls
- Target: Churn (0 = Stay, 1 = Churn)

The dataset includes 5% null values in each feature (excluding the target). Handle missing values, encode categorical variables, and scale numerical data as part of preprocessing.

3. Exploratory Data Analysis (EDA)

Use visualizations to uncover patterns and correlations:

- Countplot: Churned vs Non-Churned Customers
- Countplot: Contract Type vs Churn
- Bar Chart: Internet Service vs Churn
- Violin Plot: Monthly Charges vs Churn
- Histogram: Tenure Distribution
- Heatmap: Correlation between numerical variables
- Boxplot: Customer Support Calls vs Churn

4. Feature Engineering

Encode categorical features (One-Hot Encoding or Label Encoding), impute missing values using mean/mode, and normalize numeric columns like charges and tenure. Optionally, derive features such as 'is_long_term_customer'.

5. Model Selection

Begin with Logistic Regression for baseline results. Then explore models like Random Forest, Decision Tree, or XGBoost to improve accuracy.

6. Model Training

Split data into train-test sets (80/20). Train the selected model using features and labels.

7. Model Evaluation

Evaluate the model using the following metrics:

- Accuracy
- Precision
- Recall
- F1 Score
- Confusion Matrix
- ROC Curve

8. Model Saving and Deployment

After evaluation, save the model using `joblib` or `pickle`. Deploy it through a web service using Flask or FastAPI, with a front-end UI to allow inputs and view predictions.